





- The accuracy score represents the coefficient of determination ( $R^2$ ). This is at max 1, but can be negative. It will be 0 if you predict the mean of y for all observations.
- The R Square is measure of how close the data are to the fitted regression line.
- In this case we can say that our model explains 79% of the training data & 75% of the testing data
- The RMSE is the standard deviation of the residuals. Residuals is the difference between the predicted value and the regression line. Hence RMSE is a measure of how spread your residuals are.
- The mean absolute error (MAE) is the average of all the absolute errors. The absolute error is the difference between the true value ( $y_{train}$ ) and the predicted value ( $y_{pred}$ ).
- Coeff are the weights
- The intercept is the expected mean value of Y when all X=0

```
In [48]: # The Root Mean Squared Error (RMSE)
print('The RMSE on the training dataset is: ',sqrt(mean_squared_error(y_train,y_pred)))
print("")
print('The RMSE on the testing dataset is: ',sqrt(mean_squared_error(y_test,lm.predict(X_test))))
print("")

The RMSE on the training dataset is: 2618.561282656437

The RMSE on the testing dataset is: 2340.4200873490113
```

## Links

R2 Link: [https://scikit-learn.org/stable/modules/generated/sklearn.metrics.r2\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.r2_score.html)  
RMSE Link: [https://scikit-learn.org/stable/modules/generated/sklearn.metrics.mean\\_squared\\_error.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.mean_squared_error.html)  
MAE Link: [https://scikit-learn.org/stable/modules/generated/sklearn.metrics.mean\\_absolute\\_error.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.mean_absolute_error.html)

```
In [49]: # The Mean Absolute Error (MAE)
print('The MAE on the training dataset is: ',mean_absolute_error(y_train,y_pred))
print("")
print('The MAE on the testing dataset is: ',mean_absolute_error(y_test,lm.predict(X_test)))
print("")

The MAE on the training dataset is: 1928.067066528867

The MAE on the testing dataset is: 1692.3835577162286
```

Error is dropped for unseen data

```
In [50]: # Coefficients
print('Coefficients: ', lm.coef_)

Coefficients: [ 9.72168996e-01  3.69561477e+00 -4.29587654e+03  3.72469059e+03
 5.71185949e+02  7.38005738e+02 -3.52308895e+02 -9.02241594e+02
-1.55867567e+02  8.14821481e+02  1.50084913e+02 -2.92494076e+02]

Each of these no are corresponding to our X variable. We have 12 independent variable
```

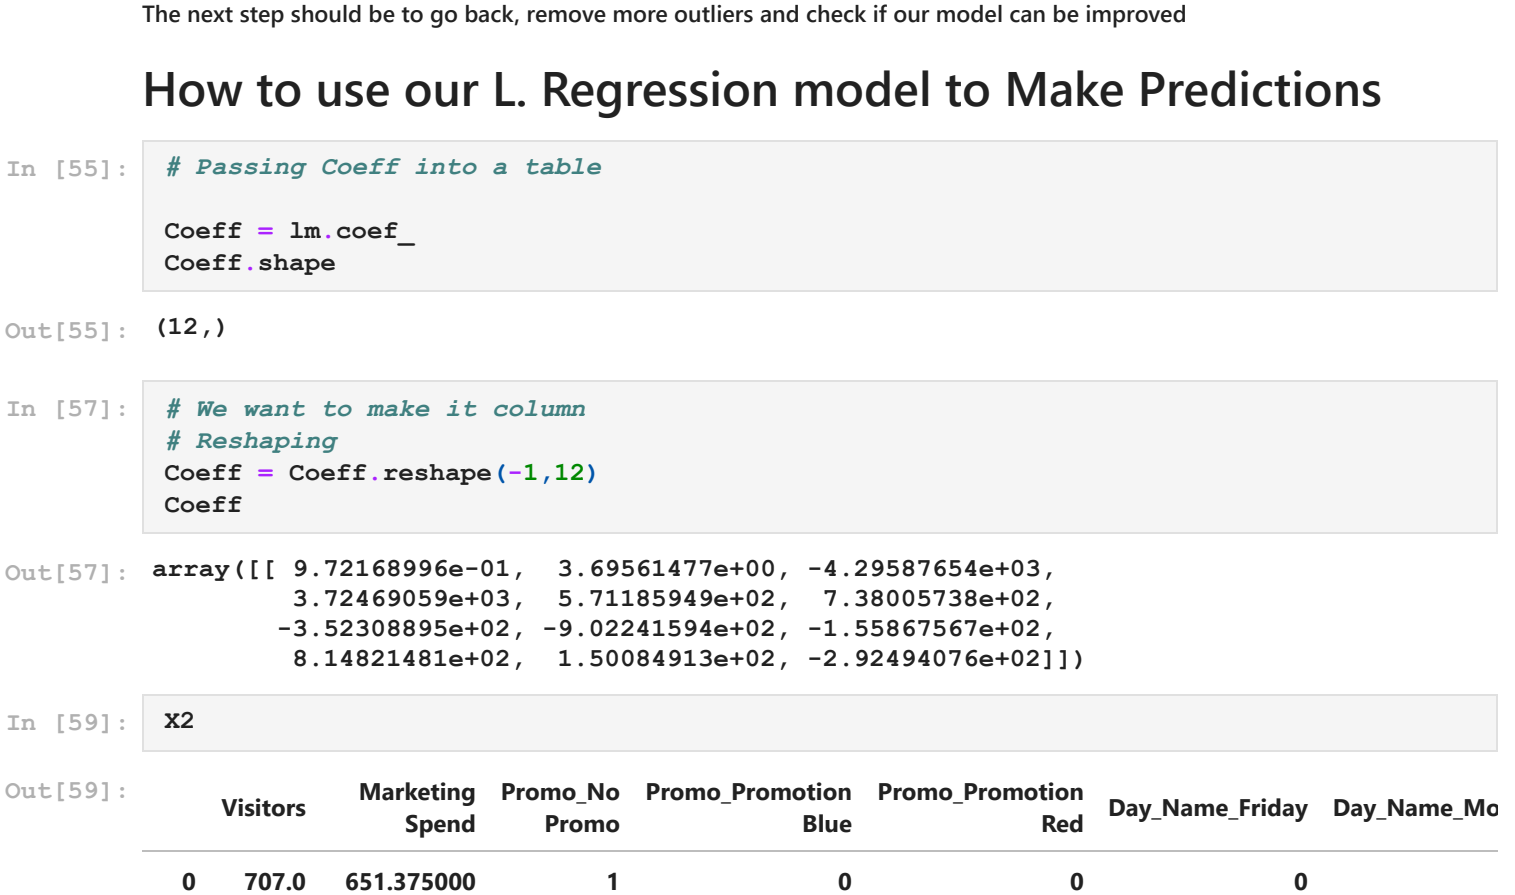
```
In [51]: # The Intercept
print('Intercept: ', lm.intercept_)

Intercept: 4060.650366499697
```

```
In [52]: # Plotting Actuals Vs Predicted

plt.figure(figsize=(15,10))

plt.scatter(y_train, y_pred, c='green')
plt.plot([y_train.min(), y_train.max()], [y_train.min(), y_train.max()], 'k--', c='red')
plt.xlabel('Actuals')
plt.ylabel('Predicted Values')
plt.title('Actuals Vs Predicted Values')
plt.grid(True)
plt.show()
```



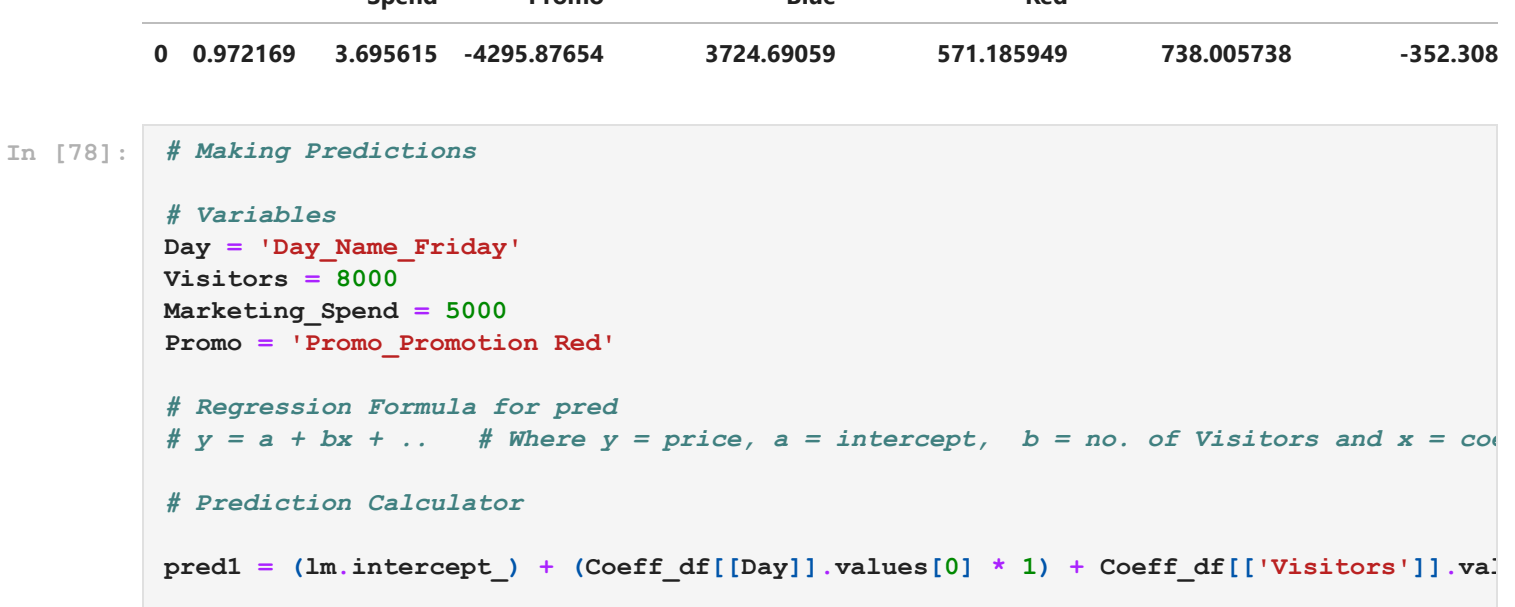
## Notes:

- This is not a regression line
- The closer data points are to the red line, better the model is.
- The further away the data points is the worst our prediction is

```
In [54]: # Plotting Residuals

plt.figure(figsize=(15,10))

sns.residplot(y_train, y_pred, color='green')
plt.xlabel('Actual Revenue')
plt.ylabel('Residuals')
plt.title('Actuals Vs Residuals')
plt.grid(True)
plt.show()
```



The next step should be to go back, remove more outliers and check if our model can be improved

## How to use our L. Regression model to Make Predictions

```
In [55]: # Passing Coeff into a table

Coeff = lm.coef_
Coeff.shape
```

Out[55]: (12,)

```
In [57]: # We want to make it column
# Reshaping
Coeff = Coeff.reshape(-1,12)
Coeff
```

Out[57]: array([[ 9.72168996e-01, 3.69561477e+00, -4.29587654e+03, 3.72469059e+03,
 -3.52308895e+02, 7.38005738e+02, -1.55867567e+02, 8.14821481e+02,
 1.50084913e+02, -2.92494076e+02]])

```
In [59]: X2

Out[59]:
```

	Visitors	Marketing Spend	Promo_No Promo	Promo_Promotion Blue	Promo_Promotion Red	Day_Name_Friday	Day_Name_Monday
0	707.0	651.375000	1	0	0	0	0
1	1455.0	1298.250000	0	0	1	0	0
2	1520.0	1559.375000	0	1	0	0	0
3	1726.0	1801.750000	1	0	0	0	0
4	2134.0	2614.500000	1	0	0	0	1
...	...	...	...	...	...	...	...
177	1400.0	1119.600000	1	0	0	0	0
178	2244.0	2067.888889	0	0	1	0	0
179	2023.0	1450.200000	1	0	0	1	1
180	1483.0	1121.875000	1	0	0	0	0
181	1303.0	871.000000	1	0	0	0	0

179 rows x 12 columns

```
In [58]: # Creating a DataFrame
Coeff_df = pd.DataFrame(Coeff, columns = ['X2_columns'])

# Displaying
Coeff_df
```

Out[58]:

	Visitors	Marketing Spend	Promo_No Promo	Promo_Promotion Blue	Promo_Promotion Red	Day_Name_Friday	Day_Name_Monday
0	0.572169	3.695615	-4295.87654	3724.69059	571.185949	738.005738	-352.308

```
In [78]: # Making Predictions

# Variables
Day = 'Day_Name_Friday'
Visitors = 8000
Marketing_Spend = 5000
Promo = 'Promo_Promotion Red'

# Regression Formula for pred
# y = a + bx + ... # Where y = price, a = intercept, b = no. of Visitors and x = co

# Prediction Calculator

pred1 = (lm.intercept_) + (Coeff_df[Day].values[0] * 1) + Coeff_df[['Visitors']].values[0]
#pred1 = (lm.intercept_) + (Coeff_df[Day].values[0] * 1) + (Coeff_df[['Visitors']].values[0] * 1)

print('The predicted Revenue is: ', pred1)

The predicted Revenue is: [31625.26784989]
```

```
In [ ]:
```